

HIGH PERFORMANCE MACHINE LEARNING ECE-GY-9143 LAB – 2

Nagharjun Mathi Mariappan
nm4074

Question C2

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 90x21
~/bin/bash
#SBATCH --account=ece-gy-9143-2022fa
#SBATCH --partition=n1s8-v100-1
#SBATCH --job-name=C222222
#SBATCH --nodes=1
#SBATCH --cpus-per-task=8
#SBATCH --output=C222222.out
#SBATCH --mem=2GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:v100:1
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision

srun python C2.py --cuda y --dataloaders 2 --opt sgd --datapath ./data
~
~
"submit_job_lab2.sh" 18L, 433C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 1.881451, Accuracy: 32.584000
Data Loading Time: 0.759291, Training Time: 8.532162, Dataloading + Training time = 9.291453, Total Time: 22.557150
-----

Epoch: 2
Loss: 1.357787, Accuracy: 50.190000
Data Loading Time: 0.721116, Training Time: 3.138121, Dataloading + Training time = 3.859237, Total Time: 16.820552
-----

Epoch: 3
Loss: 1.085082, Accuracy: 61.040000
Data Loading Time: 0.711392, Training Time: 2.760595, Dataloading + Training time = 3.471987, Total Time: 16.767476
-----

Epoch: 4
Loss: 0.882391, Accuracy: 68.390000
Data Loading Time: 0.726751, Training Time: 2.764242, Dataloading + Training time = 3.490993, Total Time: 16.754345
-----

Epoch: 5
Loss: 0.728303, Accuracy: 74.602000
Data Loading Time: 0.690628, Training Time: 2.757506, Dataloading + Training time = 3.448134, Total Time: 16.726775
-----
~
~
```

Question C3

0 Workers

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 90x21
#!/bin/bash
#SBATCH --job-name=C3_0
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C3.py --cuda y --dataloaders 0 --opt sgd --datapath ./data
~
~
~
~
~
"submit_job_lab2_0.sh" 15L, 353C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.e
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 2.127166, Accuracy: 26.182000
Data Loading Time: 14.430848, Training Time: 8.333195, Dataloading + Training Time: 22.764042, Total Time: 36.147460
-----

Epoch: 2
Loss: 1.565087, Accuracy: 42.014000
Data Loading Time: 14.405235, Training Time: 2.763377, Dataloading + Training Time: 17.168612, Total Time: 30.369579
-----

Epoch: 3
Loss: 1.279415, Accuracy: 53.556000
Data Loading Time: 14.401534, Training Time: 2.757137, Dataloading + Training Time: 17.158671, Total Time: 30.368081
-----

Epoch: 4
Loss: 1.056271, Accuracy: 62.124000
Data Loading Time: 14.384500, Training Time: 2.743560, Dataloading + Training Time: 17.128060, Total Time: 30.368712
-----

Epoch: 5
Loss: 0.908237, Accuracy: 67.826000
Data Loading Time: 14.389561, Training Time: 2.742648, Dataloading + Training Time: 17.132209, Total Time: 30.377811
-----

For 0 Workers
Total Data Loading Time: 72.011678, Total Training Time: 19.339917, Final Total Time: 157.631644
~
~
~
"C3_0.out" 51L, 2060C
```

4 Workers

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C3_4
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C3.py --cuda y --dataloaders 4 --opt sgd --datapath ./data
~
~
~
~
~
~
~
~
"submit_job_lab2_4.sh" 15L, 353C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.e
Files already downloaded and verified
Files already downloaded and verified
-----
Epoch: 1
Loss: 1.932654, Accuracy: 30.770000
Data Loading Time: 1.029289, Training Time: 5.193285, Dataloading + Training Time: 6.222574, Total Time: 18.076054
-----
Epoch: 2
Loss: 1.449043, Accuracy: 46.670000
Data Loading Time: 0.992040, Training Time: 4.239782, Dataloading + Training Time: 5.231821, Total Time: 17.182871
-----
Epoch: 3
Loss: 1.179214, Accuracy: 57.476000
Data Loading Time: 0.954112, Training Time: 4.243808, Dataloading + Training Time: 5.197921, Total Time: 17.206993
-----
Epoch: 4
Loss: 1.005200, Accuracy: 64.388000
Data Loading Time: 1.013517, Training Time: 4.259894, Dataloading + Training Time: 5.273411, Total Time: 17.270802
-----
Epoch: 5
Loss: 0.855999, Accuracy: 69.490000
Data Loading Time: 1.000116, Training Time: 4.264215, Dataloading + Training Time: 5.264331, Total Time: 17.242693
-----
For 4 Workers
Total Data Loading Time: 4.989074, Total Training Time: 22.190984, Final Total Time: 86.979414
~
~
~
"C3_4.out" 51L, 2048C
```

8 Workers

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C3_8
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C3.py --cuda y --dataloaders 8 --opt sgd --datapath ./data
~
~
~
~
~
~
~
"submit_job_lab2_8.sh" 15L, 353C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
Files already downloaded and verified
Files already downloaded and verified
-----
Epoch: 1
Loss: 1.994279, Accuracy: 28.880000
Data Loading Time: 1.339623, Training Time: 5.166684, Dataloading + Training Time: 6.506307, Total Time: 18.352231
-----
Epoch: 2
Loss: 1.524654, Accuracy: 43.800000
Data Loading Time: 1.289399, Training Time: 4.294507, Dataloading + Training Time: 5.583906, Total Time: 17.526697
-----
Epoch: 3
Loss: 1.250322, Accuracy: 54.810000
Data Loading Time: 1.295826, Training Time: 4.273381, Dataloading + Training Time: 5.569206, Total Time: 17.564771
-----
Epoch: 4
Loss: 1.040158, Accuracy: 63.074000
Data Loading Time: 1.301577, Training Time: 4.252391, Dataloading + Training Time: 5.553968, Total Time: 17.577784
-----
Epoch: 5
Loss: 0.889878, Accuracy: 68.676000
Data Loading Time: 1.301488, Training Time: 4.301344, Dataloading + Training Time: 5.602832, Total Time: 17.568652
-----
For 8 Workers
Total Data Loading Time: 6.527913, Total Training Time: 22.288306, Final Total Time: 88.590136
~
~
~
"C3_8.out" 51L, 2048C
```

12 Workers

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C3_12
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C3.py --cuda y --dataloaders 12 --opt sgd --datapath ./data
~
~
~
~
~
~
~
~
~
"submit_job_lab2_12.sh" 15L, 355C 1,1 All
```

Output:

```
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
/home/nm4074/.local/lib/python3.8/site-packages/torch/utils/data/dataloader.py:563: UserWarning: This DataLoader will create 12 worker processes
is 10, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader run
slowly/freeze if necessary.
  warnings.warn(_create_warning_msg(
Loss: 1.920384, Accuracy: 30.848000
Data Loading Time: 1.489052, Training Time: 8.656873, Dataloading + Training Time: 10.145925, Total Time: 23.801044
-----
Epoch: 2
Loss: 1.472469, Accuracy: 46.012000
Data Loading Time: 1.470722, Training Time: 2.763532, Dataloading + Training Time: 4.234254, Total Time: 17.521764
-----
Epoch: 3
Loss: 1.224550, Accuracy: 55.802000
Data Loading Time: 1.469140, Training Time: 2.749247, Dataloading + Training Time: 4.218386, Total Time: 17.512732
-----
Epoch: 4
Loss: 1.010647, Accuracy: 64.062000
Data Loading Time: 1.489399, Training Time: 2.736402, Dataloading + Training Time: 4.224801, Total Time: 17.532259
-----
Epoch: 5
Loss: 0.850490, Accuracy: 70.084000
Data Loading Time: 1.479252, Training Time: 2.756167, Dataloading + Training Time: 4.235419, Total Time: 17.524539
-----
For 12 Workers
Total Data Loading Time: 7.397564, Total Training Time: 19.661221, Final Total Time: 93.892338
"C3_12.out" 53L, 2635C
```

16 Workers

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24

#!/bin/bash
#SBATCH --job-name=C3_16
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C3.py --cuda y --dataloaders 16 --opt sgd --datapath ./data
~
~
~
~
~
~
~
"submit_job_lab2_16.sh" 15L, 355C 1,1 All
```


Output:

```
Files already downloaded and verified
Files already downloaded and verified

Epoch: 1
/home/nm4074/.local/lib/python3.8/site-packages/torch/utils/data/dataloader.py:563: UserWarning: This DataLoader will create 16 worker processes
is 10, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader slower
warnings.warn(_create_warning_msg(
Loss: 1.897020, Accuracy: 31.378000
Data Loading Time: 1.838260, Training Time: 6.216824, Dataloading + Training Time: 8.055084, Total Time: 21.878159

Epoch: 2
Loss: 1.395938, Accuracy: 48.490000
Data Loading Time: 1.861479, Training Time: 2.899351, Dataloading + Training Time: 4.760830, Total Time: 18.469837

Epoch: 3
Loss: 1.146426, Accuracy: 58.798000
Data Loading Time: 1.826175, Training Time: 2.930982, Dataloading + Training Time: 4.757157, Total Time: 18.437965

Epoch: 4
Loss: 0.970623, Accuracy: 65.192000
Data Loading Time: 1.888839, Training Time: 2.920110, Dataloading + Training Time: 4.808949, Total Time: 18.494061

Epoch: 5
Loss: 0.836521, Accuracy: 70.538000
Data Loading Time: 1.795191, Training Time: 2.826157, Dataloading + Training Time: 4.621348, Total Time: 18.363080

For 16 Workers
Total Data Loading Time: 9.209943, Total Training Time: 17.793425, Final Total Time: 95.633101
"C3_16.out" 53L, 2534C
```

20 Workers

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24

#!/bin/bash
#SBATCH --job-name=C3_20
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C3.py --cuda y --dataloaders 20 --opt sgd --datapath ./data
~
~
~
~
~
~
~
~
"submit_job_lab2_20.sh" 15L, 355C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
Files already downloaded and verified
Files already downloaded and verified
-----
Epoch: 1
/home/nm4074/.local/lib/python3.8/site-packages/torch/utils/data/dataloader.py:563: UserWarning: This DataLoader will create 28 worker processes
is 10, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader r
1 slowness/freeze if necessary.
  warnings.warn(_create_warning_msg(
Loss: 2.144719, Accuracy: 25.222000
Data Loading Time: 2.178160, Training Time: 13.448373, Dataloading + Training Time: 15.626534, Total Time: 27.956700
-----
Epoch: 2
Loss: 1.559243, Accuracy: 42.352000
Data Loading Time: 2.249210, Training Time: 4.216558, Dataloading + Training Time: 6.465768, Total Time: 18.429033
-----
Epoch: 3
Loss: 1.286485, Accuracy: 53.092000
Data Loading Time: 2.271620, Training Time: 4.253950, Dataloading + Training Time: 6.525569, Total Time: 18.491224
-----
Epoch: 4
Loss: 1.039924, Accuracy: 62.720000
Data Loading Time: 2.179353, Training Time: 4.232007, Dataloading + Training Time: 6.411360, Total Time: 18.460440
-----
Epoch: 5
Loss: 0.853286, Accuracy: 69.912000
Data Loading Time: 2.161334, Training Time: 4.195379, Dataloading + Training Time: 6.356712, Total Time: 18.481929
-----
For 20 Workers
Total Data Loading Time: 11.039677, Total Training Time: 30.346266, Final Total Time: 101.819326
"C3_20.out" 53L, 2538C
```

Inference: 4 workers are needed for optimal performance

Question C4

1 Worker

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C4_1
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C4.py --cuda y --dataloaders 1 --opt sgd --datapath ./data
~
~
~
~
~
~
~
~
~
~
"submit_job_lab2_C4_1.sh" 15L, 353C 1,1 All
```


Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 2.021689, Accuracy: 29.232000
Data Loading Time: 0.677419, Training Time: 6.634655, Dataloading + Training Time: 7.312074, Total Time: 21.183691
-----

Epoch: 2
Loss: 1.487198, Accuracy: 45.042000
Data Loading Time: 0.676466, Training Time: 2.780609, Dataloading + Training Time: 3.457075, Total Time: 16.693135
-----

Epoch: 3
Loss: 1.239478, Accuracy: 55.024000
Data Loading Time: 0.674014, Training Time: 2.772398, Dataloading + Training Time: 3.446411, Total Time: 16.703429
-----

Epoch: 4
Loss: 1.037386, Accuracy: 62.924000
Data Loading Time: 0.665432, Training Time: 2.774311, Dataloading + Training Time: 3.439744, Total Time: 16.691338
-----

Epoch: 5
Loss: 0.879116, Accuracy: 68.698000
Data Loading Time: 0.657241, Training Time: 2.767372, Dataloading + Training Time: 3.424613, Total Time: 16.686903
-----

For 1 Workers
Total Data Loading Time: 3.350571, Total Training Time: 17.729346, Total Training + dataloading time = 21.079917, Final Total Time: 87.958496
~
~
~
"C4_1.out" 51L, 2095C
```

4 Workers

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C4_4
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C4.py --cuda y --dataloaders 4 --opt sgd --datapath ./data
~
~
~
~
~
~
~
"submit_job_lab2_C4_2.sh" 15L, 353C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 1.897977, Accuracy: 31.244000
Data Loading Time: 0.938075, Training Time: 7.903067, Dataloading + Training Time: 8.841142, Total Time: 22.728416
-----

Epoch: 2
Loss: 1.438269, Accuracy: 47.286000
Data Loading Time: 0.884823, Training Time: 2.781339, Dataloading + Training Time: 3.666161, Total Time: 16.911271
-----

Epoch: 3
Loss: 1.176146, Accuracy: 58.112000
Data Loading Time: 0.872038, Training Time: 2.786107, Dataloading + Training Time: 3.658144, Total Time: 16.912045
-----

Epoch: 4
Loss: 0.980111, Accuracy: 65.232000
Data Loading Time: 0.899718, Training Time: 2.783725, Dataloading + Training Time: 3.683443, Total Time: 16.948293
-----

Epoch: 5
Loss: 0.852547, Accuracy: 70.018000
Data Loading Time: 0.913955, Training Time: 2.767763, Dataloading + Training Time: 3.681718, Total Time: 16.970042
-----

For 4 Workers
Total Data Loading Time: 4.508608, Total Training Time: 19.022001, Total Training + dataloading time = 23.530609, Final Total Time: 90.470066
~
~
~
"C4_4.out" 51L, 2095C
```

Question C5

CPU

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C5_cpu
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C5.py --cuda n --dataloaders 4 --opt sgd --datapath ./data
~
~
~
~
~
~
~
~
~
"submit_job_lab2_C5_2.sh" 15L, 355C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 1.922236, Accuracy: 31.666000
Data Loading Time: 1.000288, Training Time: 249.890459, Dataloading + Training Time: 250.890747, Total Time: 251.081424
-----

Epoch: 2
Loss: 1.451488, Accuracy: 46.588000
Data Loading Time: 0.838411, Training Time: 248.292668, Dataloading + Training Time: 249.131079, Total Time: 249.273468
-----

Epoch: 3
Loss: 1.191168, Accuracy: 57.256000
Data Loading Time: 0.847229, Training Time: 249.248548, Dataloading + Training Time: 250.095777, Total Time: 250.222515
-----

Epoch: 4
Loss: 0.986688, Accuracy: 64.614000
Data Loading Time: 0.886870, Training Time: 249.509733, Dataloading + Training Time: 250.396603, Total Time: 250.513681
-----

Epoch: 5
Loss: 0.848952, Accuracy: 70.000000
Data Loading Time: 0.863877, Training Time: 248.784109, Dataloading + Training Time: 249.647986, Total Time: 249.766862
-----

For 4 Workers
Total Data Loading Time: 4.436675, Total Training Time: 1245.725517, Total Training + dataloading time = 1250.162192, Final Total Time: 1250.857958
Average Training + dataloading time over 5 epochs = 250.032438, Average Running Time over 5 epochs: 250.171590
~
"C5_cpu.out" 53L, 2238C
```

GPU

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C5_gpu
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C5.py --cuda y --dataloaders 4 --opt sgd --datapath ./data
~
~
~
~
~
~
~
~
~
"submit_job_lab2_C5_1.sh" 15L, 355C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 1.912569, Accuracy: 31.108000
Data Loading Time: 0.906229, Training Time: 3.617638, Dataloading + Training Time: 4.523868, Total Time: 17.779190
-----

Epoch: 2
Loss: 1.462662, Accuracy: 46.364000
Data Loading Time: 0.870805, Training Time: 2.807515, Dataloading + Training Time: 3.678320, Total Time: 16.899243
-----

Epoch: 3
Loss: 1.180536, Accuracy: 57.194000
Data Loading Time: 0.883144, Training Time: 2.783385, Dataloading + Training Time: 3.666529, Total Time: 16.897940
-----

Epoch: 4
Loss: 1.005695, Accuracy: 64.060000
Data Loading Time: 0.846285, Training Time: 2.796815, Dataloading + Training Time: 3.643100, Total Time: 16.885395
-----

Epoch: 5
Loss: 0.868388, Accuracy: 69.264000
Data Loading Time: 0.872448, Training Time: 2.773811, Dataloading + Training Time: 3.646259, Total Time: 16.903017
-----

For 4 Workers
Total Data Loading Time: 4.378911, Total Training Time: 14.779165, Total Training + dataloading time = 19.158076, Final Total Time: 85.364785
Average Training + dataloading time over 5 epochs = 3.831615, Average Running Time over 5 epochs: 17.072957
~
"C5_gpu.out" 53L, 2204C
```

Question C6

Sgd

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C6_sgd
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C6.py --cuda y --dataloaders 4 --opt sgd --datapath ./data
~
~
~
~
~
~
~
~
"submit_job_lab2_C6_sgd.sh" 15L, 355C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
HERE sgd OPTIMIZER IS USED
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 2.056218, Accuracy: 26.316000
Data Loading Time: 0.934733, Training Time: 3.645197, Dataloading + Training Time: 4.579930, Total Time: 17.784543
-----

Epoch: 2
Loss: 1.464886, Accuracy: 45.820000
Data Loading Time: 0.928020, Training Time: 2.849547, Dataloading + Training Time: 3.777567, Total Time: 17.006980
-----

Epoch: 3
Loss: 1.187279, Accuracy: 57.262000
Data Loading Time: 0.917129, Training Time: 2.858217, Dataloading + Training Time: 3.775346, Total Time: 16.989703
-----

Epoch: 4
Loss: 0.998617, Accuracy: 64.628000
Data Loading Time: 0.946181, Training Time: 2.833721, Dataloading + Training Time: 3.779902, Total Time: 17.015581
-----

Epoch: 5
Loss: 0.867932, Accuracy: 69.442000
Data Loading Time: 0.942600, Training Time: 3.036116, Dataloading + Training Time: 3.978614, Total Time: 17.017943
-----

For 4 Workers
Total Data Loading Time: 4.668563, Total Training Time: 15.222797, Total Training + dataloading time = 19.891359, Final Total Time: 85.814751
~
~
"C6_sgd.out" 52L, 2122C
```

Sgdnesterov

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C6_sgdnesterov
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C6.py --cuda y --dataloaders 4 --opt sgdnesterov --datapath ./data
~
~
~
~
~
~
~
~
~
"submit_job_lab2_C6_sgdnesterov.sh" 15L, 371C 1,1 All
```


Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
HERE sgdneesterov OPTIMIZER IS USED
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 2.035512, Accuracy: 26.510000
Data Loading Time: 1.046257, Training Time: 6.019464, Dataloading + Training Time: 7.065721, Total Time: 18.750428
-----
Epoch: 2
Loss: 1.424194, Accuracy: 47.588000
Data Loading Time: 1.060095, Training Time: 4.505036, Dataloading + Training Time: 5.565130, Total Time: 17.300855
-----
Epoch: 3
Loss: 1.127429, Accuracy: 59.084000
Data Loading Time: 1.102669, Training Time: 4.526691, Dataloading + Training Time: 5.629360, Total Time: 17.353031
-----
Epoch: 4
Loss: 0.943649, Accuracy: 66.620000
Data Loading Time: 1.063306, Training Time: 4.523616, Dataloading + Training Time: 5.586923, Total Time: 17.314921
-----
Epoch: 5
Loss: 0.825601, Accuracy: 70.926000
Data Loading Time: 1.064192, Training Time: 4.473196, Dataloading + Training Time: 5.537388, Total Time: 17.296982
-----

For 4 Workers
Total Data Loading Time: 5.336519, Total Training Time: 24.048003, Total Training + dataloading time = 29.384522, Final Total Time: 88.016217
~
~
"C6_sgdneesterov.out" 52L, 2130C
```

Adadelata

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C6_adadelata
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C6.py --cuda y --dataloaders 4 --opt adadelata --datapath ./data
~
~
~
~
~
~
~
"submit_job_lab2_C6_adadelata.sh" 15L, 365C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
HERE adadelta OPTIMIZER IS USED
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 1.387872, Accuracy: 49.094000
Data Loading Time: 0.952286, Training Time: 5.140140, Dataloading + Training Time: 6.092427, Total Time: 18.701006
-----

Epoch: 2
Loss: 0.887783, Accuracy: 68.472000
Data Loading Time: 0.918600, Training Time: 5.365153, Dataloading + Training Time: 6.283753, Total Time: 18.029061
-----

Epoch: 3
Loss: 0.689228, Accuracy: 75.938000
Data Loading Time: 0.916692, Training Time: 4.347992, Dataloading + Training Time: 5.264683, Total Time: 17.993747
-----

Epoch: 4
Loss: 0.580467, Accuracy: 79.804000
Data Loading Time: 0.923176, Training Time: 4.350711, Dataloading + Training Time: 5.273887, Total Time: 18.001856
-----

Epoch: 5
Loss: 0.507765, Accuracy: 82.266000
Data Loading Time: 0.882167, Training Time: 4.348468, Dataloading + Training Time: 5.230635, Total Time: 17.962396
-----

For 4 Workers
Total Data Loading Time: 4.592920, Total Training Time: 23.552463, Total Training + dataloading time = 28.145384, Final Total Time: 90.688066
~
~
"C6_adadelta.out" 52L, 2127C
```

Adagrad

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C6_adagrad
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C6.py --cuda y --dataloaders 4 --opt adagrad --datapath ./data
~
~
~
~
~
~
~
~
~
"submit_job_lab2_C6_adagrad.sh" 15L, 363C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
HERE adagrad OPTIMIZER IS USED
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 2.170890, Accuracy: 26.312000
Data Loading Time: 1.110152, Training Time: 5.707742, Dataloading + Training Time: 6.817894, Total Time: 18.201986
-----

Epoch: 2
Loss: 1.643369, Accuracy: 38.866000
Data Loading Time: 1.059453, Training Time: 4.860693, Dataloading + Training Time: 5.920146, Total Time: 17.269872
-----

Epoch: 3
Loss: 1.377111, Accuracy: 49.688000
Data Loading Time: 1.056571, Training Time: 4.825147, Dataloading + Training Time: 5.881718, Total Time: 17.266926
-----

Epoch: 4
Loss: 1.141265, Accuracy: 59.046000
Data Loading Time: 1.046545, Training Time: 4.809437, Dataloading + Training Time: 5.855982, Total Time: 17.257288
-----

Epoch: 5
Loss: 0.978039, Accuracy: 65.068000
Data Loading Time: 1.054336, Training Time: 4.838416, Dataloading + Training Time: 5.892752, Total Time: 17.280047
-----

For 4 Workers
Total Data Loading Time: 5.327057, Total Training Time: 25.041435, Total Training + dataloading time = 30.368492, Final Total Time: 87.276119
~
~
"C6_adagrad.out" 52L, 2126C
```

Adam

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C6_adam
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C6.py --cuda y --dataloaders 4 --opt adam --datapath ./data
~
~
~
~
~
~
~
"submit_job_lab2_C6_adam.sh" 15L, 357C 1,1 All
```

Output:

```
arjun17 — nm4074@log-1:~ — ssh nm4074@gw.hpc.nyu.edu — 208x56
HERE adam OPTIMIZER IS USED
Files already downloaded and verified
Files already downloaded and verified

-----
Epoch: 1
Loss: 2.236613, Accuracy: 22.234000
Data Loading Time: 0.891450, Training Time: 5.493833, Dataloading + Training Time: 6.385284, Total Time: 19.011824
-----

Epoch: 2
Loss: 1.856090, Accuracy: 28.780000
Data Loading Time: 0.892062, Training Time: 3.938709, Dataloading + Training Time: 4.830771, Total Time: 17.509146
-----

Epoch: 3
Loss: 1.825711, Accuracy: 29.910000
Data Loading Time: 0.876408, Training Time: 3.939026, Dataloading + Training Time: 4.815434, Total Time: 17.505885
-----

Epoch: 4
Loss: 1.832697, Accuracy: 29.348000
Data Loading Time: 0.876455, Training Time: 3.927690, Dataloading + Training Time: 4.804145, Total Time: 17.519145
-----

Epoch: 5
Loss: 1.819754, Accuracy: 30.214000
Data Loading Time: 0.892791, Training Time: 3.950985, Dataloading + Training Time: 4.843776, Total Time: 17.537230
-----

For 4 Workers
Total Data Loading Time: 4.429167, Total Training Time: 21.250242, Total Training + dataloading time = 25.679409, Final Total Time: 89.083230
~
~
"C6_adam.out" 52L, 2123C
```

Question C7

```
arjun17 — nm4074@log-1:~ — ssh nm4074@greene.hpc.nyu.edu — 80x24
#!/bin/bash
#SBATCH --job-name=C7
#SBATCH --nodes=1
#SBATCH --cpus-per-task=10
#SBATCH --output=%x.out
#SBATCH --mem=64GB
#SBATCH --time=02:00:00
#SBATCH --ntasks-per-node=1
#SBATCH --gres=gpu:rtx8000
module purge
cd /scratch/nm4074/lab2

#pip3 install torch
#pip3 install torchvision
python C7.py --cuda y --dataloaders 4 --opt sgd --datapath ./data
~
~
~
~
~
~
~
"submit_job_lab2_C7.sh" 15L, 351C 1,1 All
```

Output:

```
Files already downloaded and verified
Files already downloaded and verified

Epoch: 1
Loss: 1.974142, Accuracy: 29.594000
Data Loading Time: 0.990183, Training Time: 7.064182, Dataloading + Training Time: 8.054364, Total Time: 21.340184

Epoch: 2
Loss: 1.480416, Accuracy: 45.632000
Data Loading Time: 0.959991, Training Time: 2.868290, Dataloading + Training Time: 3.828281, Total Time: 16.837974

Epoch: 3
Loss: 1.190029, Accuracy: 56.998000
Data Loading Time: 1.002573, Training Time: 2.863352, Dataloading + Training Time: 3.865925, Total Time: 16.937849

Epoch: 4
Loss: 0.981544, Accuracy: 65.168000
Data Loading Time: 0.979780, Training Time: 2.865110, Dataloading + Training Time: 3.844890, Total Time: 16.967298

Epoch: 5
Loss: 0.808178, Accuracy: 71.324000
Data Loading Time: 0.987047, Training Time: 2.866579, Dataloading + Training Time: 3.853626, Total Time: 16.983506

For 4 Workers
Total Data Loading Time: 4.919574, Total Training Time: 18.527512, Total Training + dataloading time = 23.447086, Final Total Time: 89.066811
~
~
"C7.out" 51L, 2095C
```

Q1:

How many convolutional layers are in the ResNet-18 model?

A) 18

Q2:

What is the input dimension of the last linear layer?

A) It is (512 * 1)

Q3:

How many trainable parameters and how many gradients are in the ResNet-18 model that you build (please show both the answer and the code that you use to count them) when using the SGD optimizer?

Code:

```
summary(net, (3, 32, 32))
```

Where net is where the model is loaded and (3, 32, 32) are the inputs

Number of trainable parameters = 11173962

Number of gradients = 11173962


```
arjun17 — nm4074@log-1:~ — zsh — 80x24
Conv2d-37      [-1, 512, 4, 4]      1,179,648
BatchNorm2d-38 [-1, 512, 4, 4]      1,024
Conv2d-39      [-1, 512, 4, 4]      2,359,296
BatchNorm2d-40 [-1, 512, 4, 4]      1,024
Conv2d-41      [-1, 512, 4, 4]      131,072
BatchNorm2d-42 [-1, 512, 4, 4]      1,024
BasicBlock-43  [-1, 512, 4, 4]      0
Conv2d-44      [-1, 512, 4, 4]      2,359,296
BatchNorm2d-45 [-1, 512, 4, 4]      1,024
Conv2d-46      [-1, 512, 4, 4]      2,359,296
BatchNorm2d-47 [-1, 512, 4, 4]      1,024
BasicBlock-48  [-1, 512, 4, 4]      0
Linear-49      [-1, 10]          5,130
=====
Total params: 11,173,962
Trainable params: 11,173,962
Non-trainable params: 0
-----
Input size (MB): 0.01
```

Q4:

The same question as Q3, except now use Adam (only the answer is required, not the code).

Number of trainable parameters = 11173962

Number of gradients = 11173962

```
arjun17 — nm4074@log-3:~ — ssh nm4074@gw.hpc.nyu.edu — 80x24
Conv2d-37      [-1, 512, 4, 4]      1,179,648
BatchNorm2d-38 [-1, 512, 4, 4]      1,024
Conv2d-39      [-1, 512, 4, 4]      2,359,296
BatchNorm2d-40 [-1, 512, 4, 4]      1,024
Conv2d-41      [-1, 512, 4, 4]      131,072
BatchNorm2d-42 [-1, 512, 4, 4]      1,024
BasicBlock-43  [-1, 512, 4, 4]      0
Conv2d-44      [-1, 512, 4, 4]      2,359,296
BatchNorm2d-45 [-1, 512, 4, 4]      1,024
Conv2d-46      [-1, 512, 4, 4]      2,359,296
BatchNorm2d-47 [-1, 512, 4, 4]      1,024
BasicBlock-48  [-1, 512, 4, 4]      0
Linear-49      [-1, 10]          5,130
=====
Total params: 11,173,962
Trainable params: 11,173,962
Non-trainable params: 0
-----
Input size (MB): 0.01
Forward/backward pass size (MB): 11.25
Params size (MB): 42.63
Estimated Total Size (MB): 53.89
-----
"Q3.out" 64L, 3902C                                     64,1      Bot
```